

Dataset Description

This competition provides real-world cheating detection data from a production AI interviewing system. All candidate identifiers and feature names have been anonymized to protect proprietary detection methods and user privacy.

Files

train.csv – Training set with labels and optional unlabeled examples

test.csv – Test set without labels

sample_submission.csv – Example submission file

social_graph.csv – Complete social graph as an edge list

feature_metadata.json – Feature types, value ranges, and missingness statistics

Data Format

Main Dataset Files (**train.csv**, **test.csv**, **holdout.csv**)

Each row represents a single candidate with the following columns:

user_hash – Anonymized candidate identifier

feature_001 through feature_018 – Anonymized behavioral and interview-related features.

Feature meanings are intentionally hidden to prevent reverse-engineering of cheating detection methods.

See `feature_metadata.json` for data types, value ranges, and missingness percentages.

high_conf_clean – Binary flag (1 or NaN) indicating high-confidence **not-cheating** predictions from the production model. Candidates with `high_conf_clean = 1` do not have manual review labels and represent a broader sample of the population.

is_cheating – Target variable

- 0 = Not cheating (legitimate candidate)
- 1 = Cheating (confirmed during manual review)

- NaN = Unlabeled (only for rows where `high_conf_clean = 1`)
-

Social Graph (social_graph.csv)

An edge list representing the complete social network:

user_a – Anonymized ID of one user in the connection

user_b – Anonymized ID of the other connected user

The graph includes all relationship edges in the system.

Many nodes in the graph do not appear in the main dataset (these are "ghost users").

Feature Metadata (feature_metadata.json)

Provides metadata for each anonymized:

- `type` : numeric or binary
 - `range` : min–max values or number of categories
 - `missing_%` : percentage of missing values in the training set
-

What You're Predicting

For each `user_hash` in the test set, predict the probability (`0.0` to `1.0`) that the candidate is cheating.

Important Notes

Missing Data:

Features contain realistic missingness patterns. Some checks or signals are only available for certain candidates.

NaN values are informative and should be modeled carefully.

Unlabeled Data:

The training set includes candidates

with `high_conf_clean = 1` and `is_cheating = NaN`.

These unlabeled candidates reflect a broader slice of the population.

Graph Structure:

The social graph is complete but sparse.

Not all users in the graph appear in the main dataset,
and not all users in the dataset appear in the graph.